Duckophila: a whole-brain *Drosophila* connectome model in closed loop with a bipedal robot

Technical report and preprint draft

Helge Winkler
hw@alphasix.de
https://github.com/first-it-consulting/duckophila

November 2026

Abstract

We connect the complete connectome-derived leaky integrate-and-fire model of the adult *Drosophila melanogaster* brain (FlyWire, 138,639 neurons, 54.5 million synapses; Shiu et al. 2024) to a small bipedal robot, the open-source microduck of Pollen Robotics (Hugging Face), running its production control software against a MuJoCo simulation of its body. The robot's time-of-flight depth sensor is translated into input to the fly's looming-sensitive visual projection neurons (LPLC2, LC4); its camera, through dense optic flow, into input to the fly's elementary motion detectors (T4/T5), each placed at its own visual column using FlyWire's column map. Motor commands are read out from identified descending neurons. Without any tuning of the connectome, the model (i) fires its giant fibre (DNp01) to objects approaching the robot's face (19 of 20 throws in closed loop), and (ii) produces a syndirectional optomotor response in DNp15/DNHS1 that turns the robot's head—and, through a coupling rule of ours, its walking body—with a rotating striped drum, at a closed-loop gain of 0.28 [0.24, 0.32], and (iii) turns the walking robot toward a small target through LC10a, the anterior optic tubercle and DNa02, keeping up with a target led 90° each way in 5 of 8 runs (heading gain 1.00). All three require the wiring: a degree-preserving rewired connectome with the same neurons and weights produces none (0 of 20; gain 0.00; 0 of 8). Further controls show that DNp15 is the strongest rotation-following neuron of 383 paired descending neuron types, while the escape response is only partly specific to the looming detectors and only weakly lateralised, and that the target pathway steers only for targets below the eye's equator, and more weakly to the right. To run the model in closed loop we reimplemented it as a stepped simulator that reproduces the reference Brian2 implementation spike for spike and runs faster than real time on one CPU core, using an exact closed-form bound that lets neurons unable to fire before their next input remain unsimulated. We report what the connectome contributes, what our hand-written sensory and motor interfaces contribute, and several failure modes found on the way: a dropped-input rule of the reference model, scrambled retinotopy destroying LPLC2 selectivity, the robot's own gait dominating its optic flow, and a field of view too narrow for the canonical steering neuron DNa02. Code, experiments and a recorded demonstration are open source.

1 Introduction

Complete synaptic wiring diagrams of an adult insect brain now exist. FlyWire reconstructed the whole brain of a female *Drosophila* [1], annotated every neuron with a cell type [2], and mapped the visual system to its retinotopic columns [3]. Shiu et al. [4] showed that a leaky integrate-and-fire (LIF) network built directly from this connectome—every synapse at one shared weight, its sign taken from the predicted neurotransmitter—predicts sensorimotor pathways that were then confirmed experimentally, for example from sugar-sensing gustatory neurons to the proboscis motor neuron MN9.

That model has so far been driven by prescribed stimulation and read out offline. Here we ask a different question: what does it do when its sensory neurons are driven by a robot's sensors, its descending neurons move that robot, and the consequences of those movements come back through the sensors—in closed loop, in real time? The robot is the microduck, a 15-servo bipedal robot from Pollen Robotics with an 8×8 time-of-flight (ToF) depth sensor and a head camera, controlled by reinforcement-learned locomotion policies. We do not replace the robot's locomotion: the fly brain sits above it, as a behaviour source, the way descending neurons sit above the fly's ventral nerve cord.

Our contribution is an engineering study with three parts: (1) a stepped, exact and faster-than-real-time reimplementation of the Shiu et al. model; (2) sensory interfaces that respect the fly's own anatomy—in

Table 1: What is the connectome's and what is ours. The brain model is not trained, tuned or edited.

From the connectome (unchanged)	Hand-written interfaces (ours)
Every neuron, synapse and sign; the LIF dynamics of [4]	ToF → LPLC2/LC4 Poisson rates (inverse time-to-contact)
Which neurons respond to looming, which to rotation	Camera → optic flow → T4/T5 rates by preferred direction
That the giant fibre fires (and more reliably for LPLC2/LC4 input)	The shape of the startle (sound, gaze, crouch, step back)
The sign and size of the DNp15/DNa02 left–right difference	Gains, smoothing, clamps; the body following the head
Where each T4/T5 looks (FlyWire columns); front and up, from the wiring	Approach-speed gate; efference copies of the robot's own motion
That LC10a → AOTU019/025 → DNa02 turns toward a small target, only below the equator	Target detector; bearing → LC10a rates and their dose; walking while a target is seen

particular a retinotopic placement of every T4/T5 neuron whose orientation is derived from the connectome itself; and (3) closed-loop experiments in a simulated twin of the robot that runs the robot's production daemons. Throughout, we separate what the connectome does from what our interfaces do, because that boundary is what any claim about "a fly brain controlling a robot" depends on (Table 1).

2 Methods

2.1 Brain model and a stepped reimplementation

We use the FlyWire 783 connectome as distributed with [4] (138,639 neurons, 15,091,983 connected pairs, 54,492,922 synapses) and the model's equations and constants unchanged: membrane time constant 20 ms, synaptic time constant 5 ms, rest and reset −52 mV, threshold −45 mV, refractory period 2.2 ms, synaptic delay 1.8 ms, 0.275 mV per synapse, time step 0.1 ms, and Poisson stimulation as a voltage kick of 250 × 0.275 mV.

The reference implementation in Brian2 [5] rebuilds the network for every run, which is appropriate for one-second optogenetic trials but not for a sensorimotor loop that needs tens of milliseconds of brain time per sensor frame (a 20 ms step cost 0.2 s to 0.4 s of wall time; a long run, 2.6× real time). Our stepped simulator (`flysim.py`, numba) holds the state between calls and follows Brian2's schedule within each step: state update, threshold, synaptic propagation and Poisson input, reset.

A semantic detail of the reference. In the reference model v and g are declared `(unless refractory)`. Brian2 implements this as a *conditional write*: any write to these variables, including synaptic input, is discarded while the neuron is refractory. Input arriving during refractoriness is therefore lost, not accumulated. An implementation that banks it is plausible-looking but wrong: every downstream neuron fired ∼20% too fast in ours until this was matched. We found it by replaying the reference's own input spike times into our simulator and comparing every downstream spike.

Exact sparsity. Under strong visual drive ∼50,000 neurons sit just off rest, and stepping them made the brain four times slower than real time; parallelising the per-step pass across cores made it slower still (thread-region launch every 0.1 ms). Between inputs the model is linear, so the highest a neuron's voltage can still reach without further input has a closed form:

$$v_{\max} - v_0 \le \max(u, 0) + P_g \max(g, 0), \qquad P_g = \max_t \frac{\tau}{\tau - T}\left(e^{-t/\tau} - e^{-t/T}\right) \approx 0.157, \tag{1}$$

with $u = v - v_0$, T the membrane and τ the synaptic time constant. A non-refractory neuron whose bound is below threshold cannot spike before its next input; it becomes *dormant* and is not stepped. Its next synaptic arrival or Poisson event wakes it, and its state is advanced analytically over the gap (the per-step propagator raised to the gap, from a lookup table). Poisson input is drawn as geometric inter-event gaps per frame rather than a Bernoulli trial per neuron per step, which is the same process. The result remains spike-for-spike identical to Brian2 under replay (13,607 of 13,607 spikes) while cost follows synaptic events instead of neurons × steps.

2.2 Robot, simulator and interface

The microduck's production daemons (`robotd`, `tofd`, `mediad`) run unmodified against a MuJoCo [6] model of its body, with only the servo bus, the ToF sensor and the camera replaced by the simulator (`scripts/duck-sim`). Our bridge reads the robot's own interfaces: ToF frames at 15 Hz (`tof.stream`), camera frames at 30 Hz (`media.frame`), and the robot's state at 50 Hz (`robot.state`: head forward kinematics, IMU orientation, odometry, executed twist). It writes only intents the robot already accepts—`robot.move`, `robot.head`, `robot.look`, `robot.pose`, `robot.sound`. The brain advances by exactly the time between ToF frames, with the latest camera-derived rates.

A stimulus prop (a sphere that can be thrown at the robot) and a rotating drum of 48 alternating black and white panels (0.5 m radius, a 15° spatial period) are added to the simulator without modifying it: a generated copy of the scene, and a Python start-up hook that moves the two mocap bodies.

The twin had to walk first. The twin as shipped loaded an obsolete walking policy on a robot model without joint backlash, and did not walk: 4 cm in 8 s at 0.15 m/s commanded. The current policy (velstand) was trained with backlash and stands still without it; with both corrected the twin walks (0.06 m/s for 0.15 m/s asked) and turns (50 ř/s for 46 ř/s asked on the spot). This fix has been submitted upstream. The residual shortfall in forward speed, a dead band below ∼0.55 rad/s of yaw on the spot, and an occasional failure to take a first slow step are properties of the twin that our controllers have to live with.

2.3 Startle: depth sensor to looming detectors

The giant fibre (DNp01, one per side) receives its two largest excitatory inputs from LPLC2 and LC4 (10.6% and 8.6% of the left GF's annotated input synapses; 12.1% and 8.4% of the right's), consistent with their established roles in looming detection [9, 10, 8]. For each ToF zone we estimate inverse time-to-contact $\tau^{-1} = -\dot{d}/d$ from a line fit over four frames (0.27 s); a zone drives the LPLC2 and LC4 neurons assigned to it (left four columns to the left optic lobe, round-robin within a side, ∼3 LPLC2 and ∼1.5 LC4 per zone) at a saturating rate up to 150 Hz. Two filters are ours and labelled as such: zones whose return is the floor are excluded, using the robot's own floor test with a 3° pose margin; and drive is gated on approach speed above 0.4 m/s, a duck choice that loosely parallels the giant fibre's preference for fast looming [9]. A giant-fibre spike triggers a startle assembled from existing intents, with a 2 s refractory period.

2.4 Optomotor reflex: camera to motion detectors

Retinotopy from the connectome. Each eye has ∼750 T4 and ∼750 T5 neurons of each of four preferred directions (a front-to-back, b back-to-front, c up, d down [7]). FlyWire's column map [3] assigns 11,822 of them to 796 hexagonal columns (395 left, 401 right) with axial coordinates (p, q). Nearest neighbours differ by $(p{+}1, q{+}1)$ or by one coordinate, which gives physical coordinates $X = \sqrt{3}\,(q - p)/2$, $Y = (p + q)/2$ and a roughly round eye.

The map does not say which direction is anterior or dorsal. We derive both from the wiring (Fig. 3). LPLC2's four dendritic arms each take the T4/T5 subtype whose preferred direction points away from the cell's centre [8]; averaging, over all LPLC2 of an eye, the position of each subtype's inputs relative to the cell's input centroid places T4a/T5a at $\Delta X = -2.7$ and -2.9 columns (left and right eye), T4b/T5b at $+2.1$ and $+1.9$, T4c/T5c at $\Delta Y = +1.7$ and $+1.9$, and T4d/T5d at -2.1 and -2.2. Hence $+X$ is anterior and $+Y$ dorsal in both eyes. Independently, the dorsal and ventral HS cells (HSN, HSS) take their T4/T5 input at higher and lower Y respectively, confirming the vertical axis, and the equatorial HSE sets elevation zero. With an interommatidial angle of 4.8° and the frontmost column placed 10° across the midline, the eyes span $-10°$ to 131° in azimuth and $-84°$ to 66° in elevation.

Flow to rates. The camera sees 45°×73° straight ahead. Only T4/T5 neurons whose direction falls in that view—1,424 of 11,822—receive input: dense Farnebäck optic flow [15] at quarter resolution, averaged over one column's extent at the neuron's position, projected on its preferred direction, rectified and mapped to a saturating Poisson rate (≤60 Hz). Image left is the robot's left (verified with a placed object).

Efference copies. Image motion the robot causes deliberately is predicted and subtracted before it reaches the fly: the head's idle look-around, and walking. Walking turned out to be dominated not by translation (floor flow ∼10 ř/s at 0.08 m/s) but by the gait's pitch and roll (30 ř/s to 80 ř/s). We predict both per pixel from the robot's state—floor translation from odometry on pixels whose ray meets the floor, rotation from the camera's world orientation between states with the vertical component removed (yaw slip is the reflex's own signal)—and subtract the prediction weighted by local image texture, since a flow estimator, like a

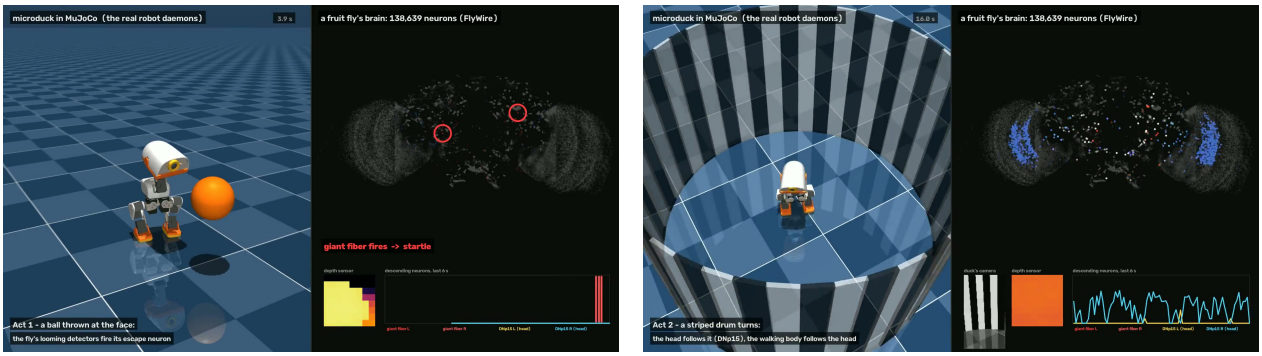


Figure 1: Frames from the recorded demonstration (`media/duckophila.mp4`). Left halves: the robot in MuJoCo, re-rendered from its recorded pose. Right halves: the brain, every neuron at its FlyWire position seen from the front, lit in its class's colour when it spikes. **Left:** a ball reaches the robot's face; both giant fibres are ringed in red as they fire. **Right:** in the turning drum, T4/T5 neurons in both optic lobes (blue) and DNp15 respond.

T4 neuron, reports no motion on an edgeless surface. The prediction matches a controlled head nod with $r = 0.98$; during walking it correlates with the measured flow ($r = 0.90$) only when each frame is paired with the state at its own capture time.

Motor readout. DNp15 (DNHS1 [11]) left–right difference, smoothed (0.15 s), integrates into head yaw. For the body we tested DNa02 [12] directly (`-steer dna02`) and a head-led coupling of ours (`-steer body`): the walking body (0.1 m/s) turns at $k\cdot$ head yaw while the head counter-rotates by the trunk's measured yaw rate, so gaze is held.

2.5 Following: a camera target to the small-object pathway

FlyWire's brain is female, so it lacks P1 and the courtship-pursuit circuit. It does contain the visual object pathway that pursuit uses: LC10a, lobula columnar neurons tuned to small moving objects [18, 19], project to the anterior optic tubercle, whence AOTU019 (GABAergic, mostly onto the contralateral DNa02) and AOTU025 (cholinergic) reach DNa02. We locate each LC10a's receptive field as the synapse-weighted mean viewing direction of its inputs that lie on the column map (Tm20, Tm4 and others), on its own eye: 232 of 234 LC10a receive such input. The target—in simulation a small self-lit sphere, the only saturated object in the grey camera image, found as the largest bright blob after a morphological opening—gives a bearing and elevation through the pinhole model, and drives each LC10a at $r_{\max}\exp(-d^2/2\sigma^2)$, d the angle between target and receptive field ($r_{\max} = 120$ Hz, $\sigma = 10°$, chosen in a dose sweep in which DNa02 is graded; at 250 Hz it saturates and folds). The readout is DNa02's left–right difference, smoothed (0.4 s), as yaw rate while walking at 0.1 m/s; with no target in view the robot stands, and the head is held straight so that the camera faces where the body goes. In closed loop, a target on the floor 0.45 m ahead of the robot swings round it at 8 ř/s, 90° left, back, 90° right and back (65 s), scored from the simulator's ground truth.

3 Results

3.1 The stepped model reproduces the reference

With the sugar-sensing gustatory neurons driven at 150 Hz for 30 one-second trials, firing rates of the 450 active neurons correlate with the Brian2 reference at $r = 0.9998$; MN9 fires at 83.2 Hz against 83.9 Hz (Fig. 2). Under replay of the reference's own input spikes, all 13,607 downstream spikes are identical in neuron and time step. On one core of an Apple M5 Max, one second of brain takes 0.04 s under that stimulus, 0.38 s under drum-like drive of 3,000 T4/T5 at 30 Hz, and 0.78 s with all 12,246 T4/T5 at 30 Hz; in the closed loop, a 67 ms frame costs 4 ms to 7 ms.

3.2 Startle

Driving LPLC2 and LC4 of one side directly, the giant fibres are silent at rest and fire in a graded way with drive rate and with the fraction of the population driven, more strongly ipsilaterally; at higher drive the contralateral DNa01/DNa02 and, weakly, the moonwalker descending neurons (MDN [13]) are recruited. On synthetic ToF scenes ray-cast with the sensor's beam geometry and noise (a 5 cm ball approaching from 1.2 m, left or right, 40 trials per speed), the giant fibre fired in 0/40 trials at 0.2 m/s, 21/40 at 0.4 m/s, 40/40

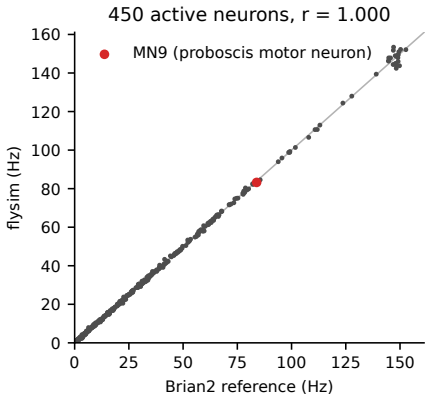


Figure 2: The stepped simulator against the Brian2 reference: mean firing rates over 30 trials with the right sugar-sensing neurons at 150 Hz (FlyWire 630, as in [4]).

at each of 0.6 m/s to 1.5 m/s, and 28/40 at 2.0 m/s, with a mean latency of 0.54 s from onset for approaches ≥1 m/s. Through the full robot stack in the simulator, 19 of 20 balls at 1.5 m/s produced a startle (0.95, 95% CI [0.76, 0.99]), latency 0.53 s, and 0 of 20 hands at 0.2 m/s (Table 2). The speed tuning is largely ours: the approach-speed gate imposes the cut, and the input null below shows the same cut. Very fast approaches are missed 30% of the time by the encoder, and walking towards a wall at 0.3 m/s produced a false startle in 6 of 20 trials.

Two failure modes were found and fixed in the interface, not the brain. A standing robot's idle head motion makes ∼45 of 64 depth zones see an approaching and receding floor; without a floor filter 30% of quiet frames drove the looming detectors. Near the horizon, a 2° to 3° discrepancy between the kinematic model and the rendered sensor pose decides whether a grazing return is floor; without a pose margin, the robot startled 5 times in 25 s in an empty room, with it none.

3.3 Optomotor reflex

In the connectome. Driving T4a/T5a of one eye and T4b/T5b of the other—the pattern a yaw rotation of the world produces—activates the ipsilateral HS cells and contralateral H2 as in the fly, and flips the left–right activity of DNa02 with the direction of rotation, in the sense of a syndirectional (following) turn under the interpretation of [12]: world turning right, DNa02 right 28 Hz versus left 5 Hz; world turning left, left 69 Hz versus right 5 Hz. DNa01 does not flip cleanly. The strongest-flipping descending neuron is DNp15 (DNHS1).

Retinotopy is necessary. With T4/T5 neurons assigned to image regions at random, wide-field drum motion reached LPLC2 scrambled and fired the giant fibre 7–11 times per run. With column-based placement, rotation produced no giant-fibre activity in synthetic tests or in the closed loop (one startle in 29 s of drum, from the ToF at start-up).

The camera's field of view limits the readout. With the column map, a whole-eye yaw rotation flips both DNa02 and DNp15 (Fig. 4). Restricted to the neurons whose direction falls in the camera's view, DNp15 still flips strongly (51 Hz left vs. 0 Hz right; 0 vs. 86 Hz), but DNa02 does not change sign: right exceeds left in both directions, by 2.7 Hz and 11 Hz. Rotation over the frontal 30° of the eye at all elevations is sufficient for a DNa02 flip; the camera's limited elevation range is not. Likewise, no expansion stimulus within the camera view, up to 400 ř/s, drove LPLC2 appreciably; looming therefore remains the ToF's task.

Closed loop. With DNp15 driving head yaw, the robot's head follows the drum in both directions as a triangle wave at 8 ř/s to 12 ř/s for a 30 ř/s drum (closed-loop gain 0.28, 95% CI [0.24, 0.32], over 12 drum phases: partial compensation) and holds still when the drum is still; higher gains oscillate, given ∼0.3 s of loop latency. With the head-led body coupling, the walking robot turns with the drum at a mean of +5 and −9 ř/s (peaks ±15 ř/s to 20 ř/s) against an intrinsic gait drift of −4.6 ř/s. Without the walking efference copy, three startles occurred in 200 s of walking in the turning drum; with it, none, with unchanged steering.

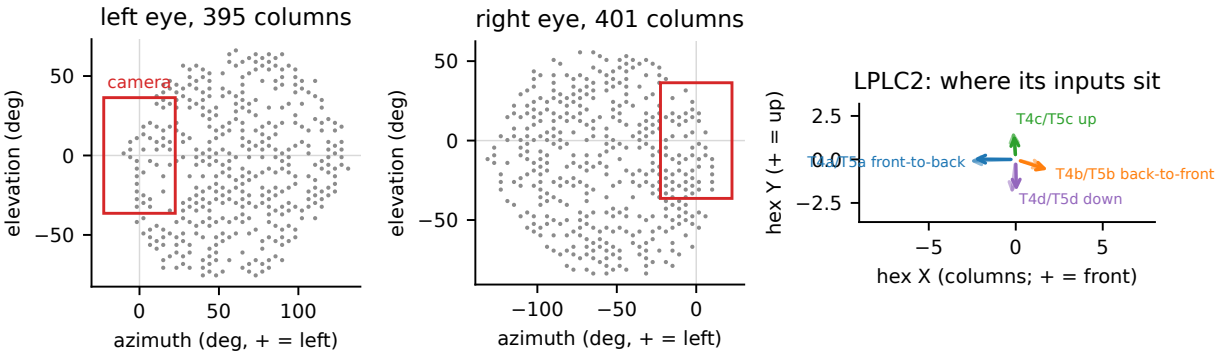


Figure 3: Retinotopy from the connectome. **Left, centre:** viewing directions of the 395 and 401 columns of each eye, from FlyWire's column map with orientation derived from the wiring; red, the camera's field of view. **Right:** mean offset of each T4/T5 subtype's input to LPLC2 from the LPLC2 neuron's input centroid (solid: left eye; faint: right). Front-to-back inputs sit posterior, up inputs dorsal—the outward-motion arrangement of a looming detector—which fixes the eye's axes.

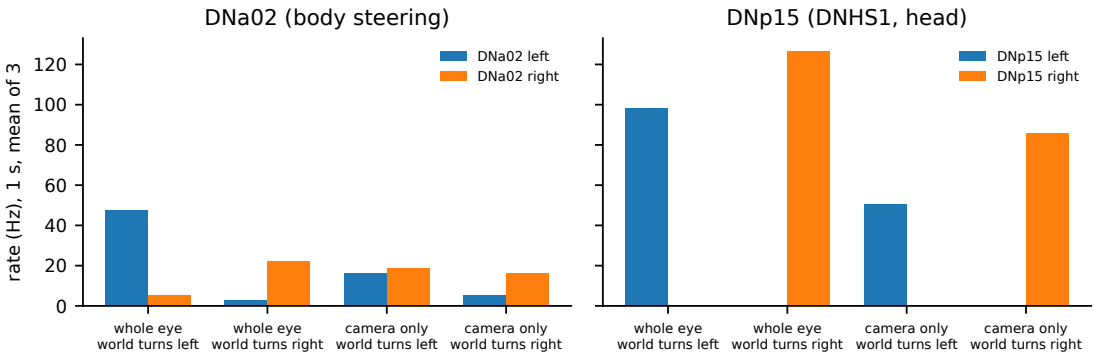


Figure 4: Steering descending neurons under yaw rotation (40 ř/s), with every T4/T5 driven (whole eye) or only those looking into the camera's view. DNp15 flips in both cases; DNa02 only with the whole eye.

3.4 Following

Open loop. A target 10° below the eye's equator makes DNa02 turn toward it, graded with bearing: left minus right −5.5, −1.8, +1.6, +5.3 and 14.2 Hz at −20, −10, 0, 10 and 20° (+ = left; 20 seeds; Fig. 5). Across 383 paired descending neuron types the left–right difference of nine tracks the bearing ($r > 0.5$); DNa02 ranks third ($r = 0.94$), after DNae002 and DNa16. The rewired brain gives no DNa02 response, and with the same rates sent to as many random other visual projection neurons the bearing dependence vanishes ($r = 0.05$). Two properties of the pathway shape the behaviour. It responds to targets below the equator only: the turn toward a target at ±10° is 0 Hz at 5° elevation and above, 2.1 Hz at −5° and 3.9 Hz to 4.4 Hz from −10° down; a target at the camera's height, which appears above the image centre, is ignored, so the target is on the floor. And it is asymmetric: the left eye's LC10a receptive fields lie further anterior than the right's (a target dead ahead drives the left population at 931 Hz in total against 782 Hz), so DNa02 leans left and a target to the right gets about a third of the turn of one to the left. Whether this asymmetry is the fly's or our column-to-angle approximation we cannot tell.

Closed loop. In 5 of 8 runs the robot kept the target in view throughout, with a heading gain of 1.00 and a median error of 7.8° [6.6, 9.3]. In the other three it lost the target where the target reversed from the leftward swing into a rightward turn, fell more than the camera's half-width behind and, with nothing in view, stood. The rewired brain never followed (0 of 8, 95% CI [0.00, 0.32]; gain −0.01). No run of either brain produced a startle.

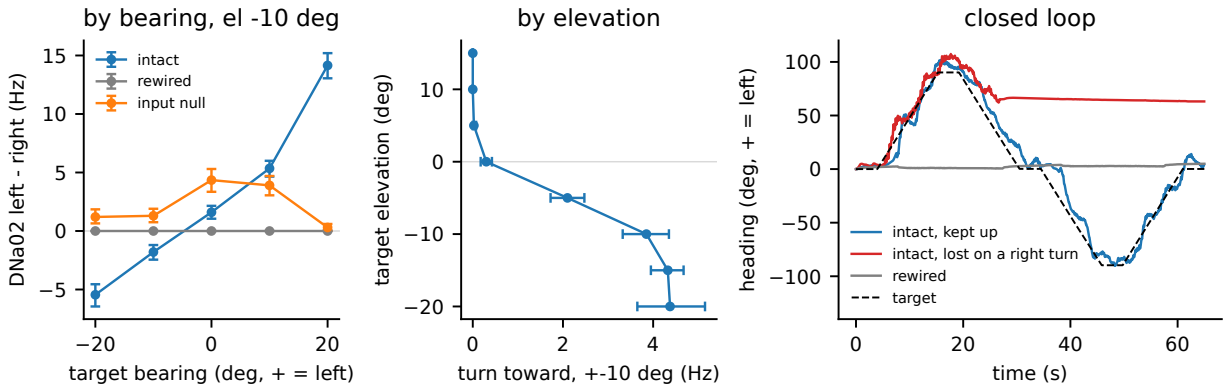


Figure 5: Following. **Left:** DNa02 left minus right against the target's bearing, open loop (mean, 95% CI, 20 seeds), for the intact brain, the rewired brain, and the input null. **Centre:** the turn toward a target 10° out, against its elevation. **Right:** closed loop, the robot's heading against the target's for a run that kept up, one that lost the target on a right turn, and the rewired brain.

Table 2: Controls. Open loop: 20 seeds per condition (40 trials per speed across both sides); closed loop: 20 throws, 8 following runs, through the full robot stack. 95% Wilson (proportions) or bootstrap (means) intervals.

	intact	rewired	sign-flipped	input null
Startle, 0.6 m/s to 1.5 m/s, open loop	1.00 [0.91, 1.00]	0.00	0.00	0.50 [0.35, 0.65]
Startle, 0.2 m/s, open loop	0.00 [0.00, 0.09]	0.00	0.00	0.00
Startle, 1.5 m/s, closed loop	0.95 [0.76, 0.99]	0.00 [0.00, 0.16]		
Startle, 0.2 m/s, closed loop	0.00 [0.00, 0.16]	0.00 [0.00, 0.16]		
Ipsilateral GF spike share	0.58 [0.56, 0.59]	–	–	0.54 [0.52, 0.56]
False startle, walking at a wall 0.3 m/s	6/20	0/20	0/20	2/20
Head optomotor gain, closed loop	0.28 [0.24, 0.32]	0.00		
DNp15 follows rotation (camera view)	yes, rank 1/383	no	no	
DNa02 follows rotation	whole eye only	no	no	
DNa02 tracks target bearing, open loop	$r = 0.94$, rank 3/383	$r = 0.00$		$r = 0.05$
Following, runs kept up, closed loop	5/8 [0.31, 0.86]	0/8 [0.00, 0.32]		

3.5 Controls

To ask what the connectome contributes, we repeated each experiment with brains that are wrong in one respect (Table 2): a *rewired* connectome, in which every connection keeps its presynaptic neuron and weight but receives a target from a random permutation of all targets (so every neuron keeps its in- and out-degree); a *sign-flipped* connectome; an *input null*, the intact brain with the looming encoder's output sent to as many randomly chosen other visual projection neurons of the same side; and a *readout null*, every descending neuron type with exactly one neuron per side scored for following rotation.

The rewired and sign-flipped brains remain active—55 000 spikes per second under visual drive, against 66 000 intact—but activity no longer reaches the descending neurons: no startle in any of 280 open-loop or 20 closed-loop trials, no DNp15 or DNa02 response, and a closed-loop head gain of 0.00. DNp15's following is specific: of 383 paired descending neuron types, 13 follow a rotation within the camera's view (95% CIs excluding zero in both directions) and DNp15 follows most strongly (rank 1, as also over the whole eye). The escape response is less specific than the demonstrations suggest: the input null still fired the giant fibre in about half of fast approaches (0.50 [0.35, 0.65] at 0.6 m/s to 1.5 m/s, against 1.00 for LPLC2/LC4), so in this model the giant fibre can be reached from visual input in general, and LPLC2/LC4 make it reliable. Laterality is weak: the giant fibre ipsilateral to the ball carried 58% of giant-fibre spikes (95% CI [56, 59]) in open loop and 57% [45, 66] in closed loop, close to the input null's 54% [52, 56].

4 Discussion

The connectome, unmodified and untrained, produced three recognisable fly behaviours when its sensory neurons were driven by a robot's sensors: an escape decision to fast approaching objects, a following optomotor response, and turning toward a small object. None was wired in, and all disappear when the same neurons and weights are rewired. Their specificity differs: the optomotor readout singles out DNp15 among hundreds

of candidates, whereas the escape decision is made more reliable, but not made possible, by the looming detectors, and carries little information about the side of the threat. An earlier, demonstration-based version of these results described the escape as lateralised; the controls do not support that. What was wired in is listed in Table 1, and most of the effort went there: the biological plausibility of the input decided whether the connectome's computation could be expressed at all. Scrambled retinotopy did not just add noise, it produced false escape decisions, because LPLC2's selectivity is a property of where its inputs sit. The robot's own movement dominated its visual input, as it does for animals, and needed an efference copy before the fly's motion pathway could be used while walking [14].

The clearest limitation is sensory. A single forward camera covers a small part of a fly's visual field, and the canonical steering neuron DNa02 [12] did not respond to rotation within it. The head-coupled neuron DNp15 did, so we steer the head with the fly's output and the body with a coupling rule of ours. Whether a wider camera, or several, recover DNa02-based steering is a direct experimental question the setup can answer. For a small target DNa02 does respond within the camera's view, so following is steered by the fly's own steering neuron; but its losses—a target that leaves the 45° view is not found again—are again the camera's, against a fly's nearly panoramic eyes.

Other limitations are those of the model and the platform. The LIF model of [4] uses one synaptic weight for all connections, has no gap junctions (the giant fibre's electrical synapses are absent), no plasticity and no neuromodulation, and treats neurotransmitter predictions as signs. T4 and T5 are driven alike, without an ON/OFF split. The robot is simulated; its gait reaches 40% of commanded speed and the camera/state timing measured in simulation will differ on hardware. Our gains and gates are engineering choices, stated as such. LC10a's preference for moving over stationary objects is not modelled. We make no claim that the robot behaves like a fly beyond the three behaviours measured.

Natural next steps are valence learning in the mushroom body with plasticity at Kenyon-cell output synapses, and the physical robot, with the brain on a laptop and the robot's camera streamed over WebRTC. Related efforts embody connectome-constrained or neuromechanical fly models in simulated bodies [16, 17]; here the body is a different animal's, and the interesting part is the translation between them.

Reproducibility

Everything above is a script in `https://github.com/first-it-consulting/duckophila`: `experiments/` (each figure and number), `paper/figures.py` (every figure, from data), `demo.sh -video` (the recorded demonstration). The FlyWire data are fetched under their own terms. Hardware: Apple M5 Max, macOS; brain on one core.

Acknowledgements

The connectome and annotations are the work of the FlyWire consortium; the brain model is that of Shiu et al.; the robot, its software and its simulator are Pollen Robotics'. The FlyWire annotations used are version $\geq$3.2, for which the annotation authors also ask to cite Berg et al. (2026) and Tastekin et al. (2026). The implementation, experiments and this report were developed with substantial assistance from an AI coding assistant (Anthropic's Claude), directed and reviewed by the author.

References

[1] S. Dorkenwald et al. Neuronal wiring diagram of an adult brain. *Nature* 634, 124–138 (2024).

[2] P. Schlegel et al. Whole-brain annotation and multi-connectome cell typing of *Drosophila*. *Nature* 634, 139–152 (2024).

[3] A. Matsliah et al. Neuronal parts list and wiring diagram for a visual system. *Nature* 634, 166–180 (2024).

[4] P. K. Shiu et al. A *Drosophila* computational brain model reveals sensorimotor processing. *Nature* 634, 210–219 (2024).

[5] M. Stimberg, R. Brette, D. F. M. Goodman. Brian 2, an intuitive and efficient neural simulator. *eLife* 8, e47314 (2019).

[6] E. Todorov, T. Erez, Y. Tassa. MuJoCo: a physics engine for model-based control. *IROS* (2012).

[7] M. S. Maisak et al. A directional tuning map of *Drosophila* elementary motion detectors. *Nature* 500, 212–216 (2013).

[8] N. C. Klapoetke et al. Ultra-selective looming detection from radial motion opponency. *Nature* 551, 237–241 (2017).

[9] C. R. von Reyn et al. Feature integration drives probabilistic behavior in the *Drosophila* escape response. *Neuron* 94, 1190–1204 (2017).

[10] J. M. Ache et al. Neural basis for looming size and velocity encoding in the *Drosophila* giant fiber escape pathway. *Curr. Biol.* 29, 1073–1081 (2019).

[11] S. Namiki, M. H. Dickinson, A. M. Wong, W. Korff, G. M. Card. The functional organization of descending sensory-motor pathways in *Drosophila*. *eLife* 7, e34272 (2018).

[12] A. Rayshubskiy et al. Neural circuit mechanisms for steering control in walking *Drosophila*. *bioRxiv* (2020).

[13] S. S. Bidaye, C. Machacek, Y. Wu, B. J. Dickson. Neuronal control of *Drosophila* walking direction. *Science* 344, 97–101 (2014).

[14] A. J. Kim, J. K. Fitzgerald, G. Maimon. Cellular evidence for efference copy in *Drosophila* visuomotor processing. *Nat. Neurosci.* 18, 1247–1255 (2015).

[15] G. Farnebäck. Two-frame motion estimation based on polynomial expansion. *Scandinavian Conference on Image Analysis* (2003).

[16] V. Lobato-Rios et al. NeuroMechFly, a neuromechanical model of adult *Drosophila melanogaster*. *Nat. Methods* 19, 620–627 (2022).

[17] R. Vaxenburg et al. Whole-body physics simulation of fruit fly locomotion. *Nature* (2025).

[18] I. M. A. Ribeiro et al. Visual projection neurons mediating directed courtship in *Drosophila*. *Cell* 174, 607–621 (2018).

[19] T. Hindmarsh Sten, R. Li, A. Otopalik, V. Ruta. Sexual arousal gates visual processing during *Drosophila* courtship. *Nature* 595, 549–553 (2021).